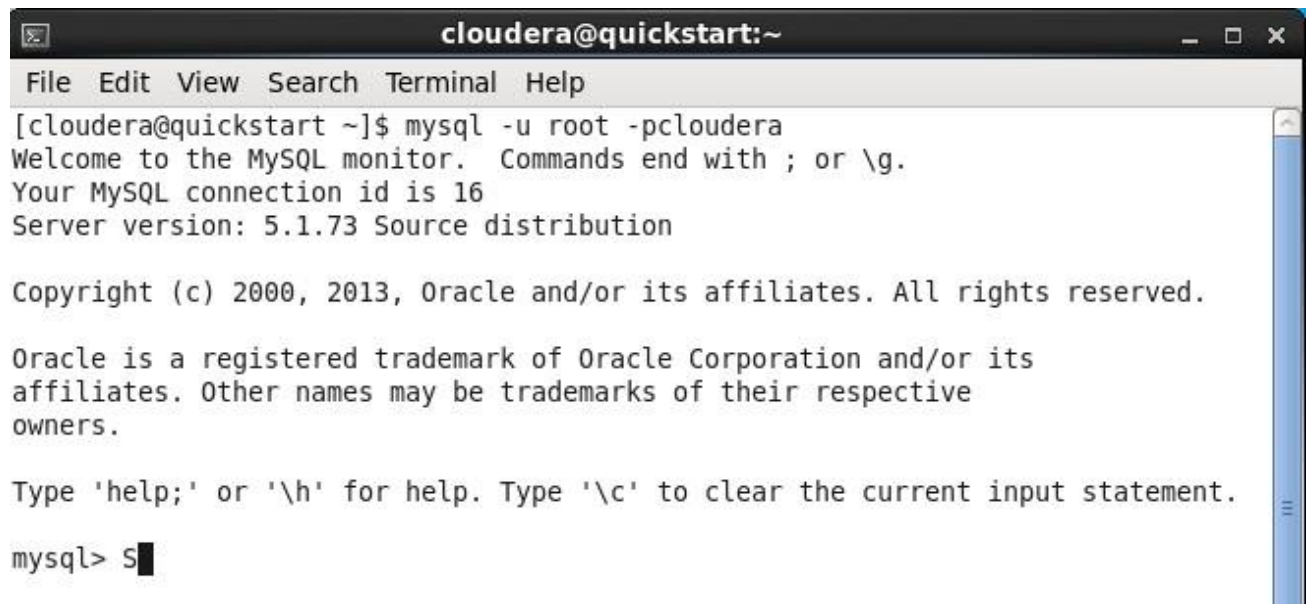


## EXPERIMENT NO. 2

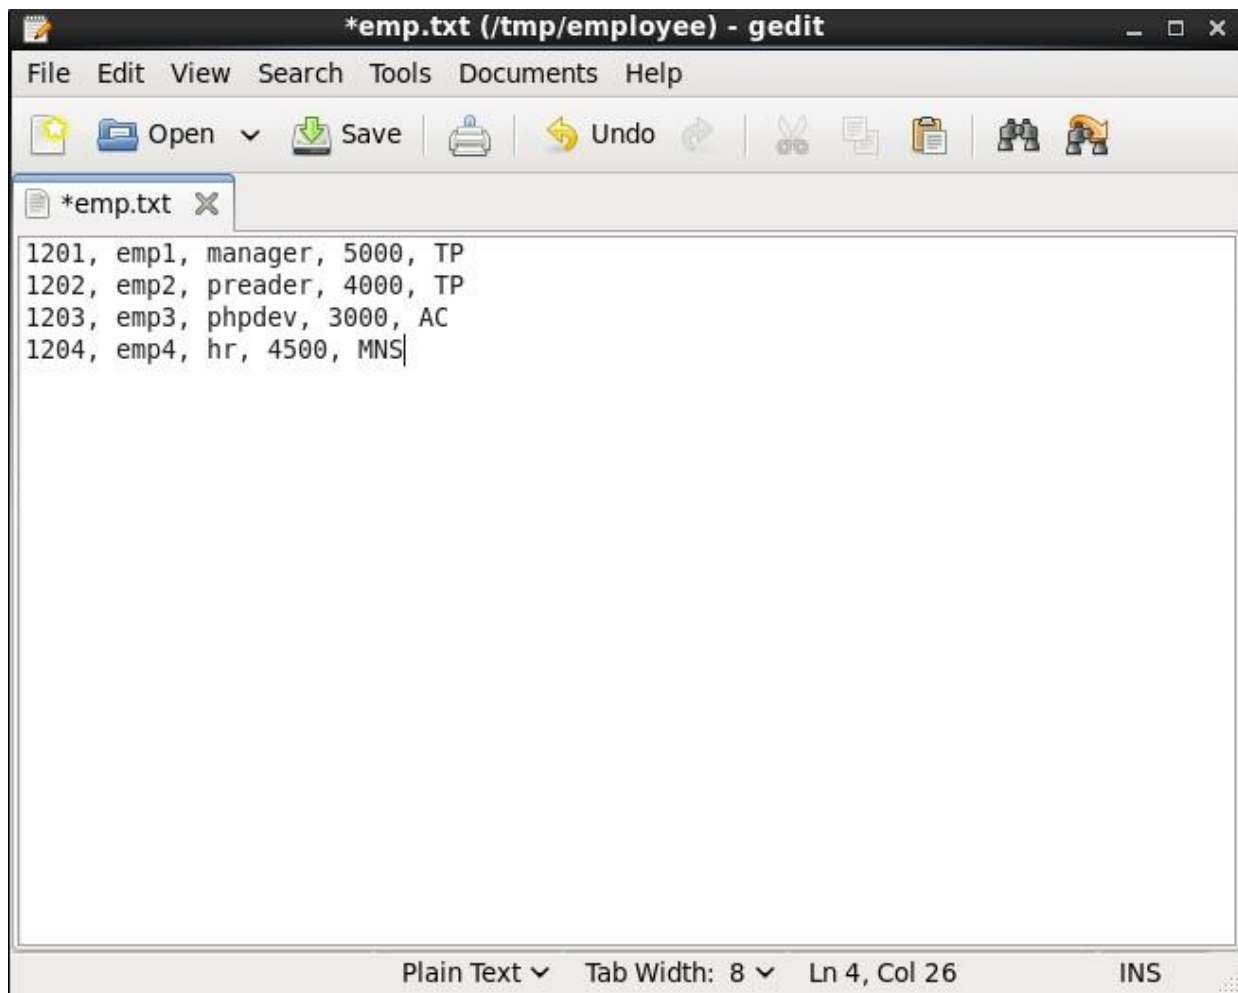
\$mysql -u root -pcloudera

A terminal window titled 'cloudera@quickstart:~' with a menu bar (File, Edit, View, Search, Terminal, Help). The terminal shows the command 'mysql -u root -pcloudera' being executed. The output includes a welcome message, connection ID 16, server version 5.1.73, and copyright information. The prompt 'mysql>' is followed by a space and a cursor.

```
cloudera@quickstart:~  
File Edit View Search Terminal Help  
[cloudera@quickstart ~]$ mysql -u root -pcloudera  
Welcome to the MySQL monitor.  Commands end with ; or \g.  
Your MySQL connection id is 16  
Server version: 5.1.73 Source distribution  
  
Copyright (c) 2000, 2013, Oracle and/or its affiliates. All rights reserved.  
  
Oracle is a registered trademark of Oracle Corporation and/or its  
affiliates. Other names may be trademarks of their respective  
owners.  
  
Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.  
mysql> S
```

```
cloudera@quickstart:~  
File Edit View Search Terminal Help  
owners.  
  
Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.  
  
mysql> show databases;  
+-----+  
| Database |  
+-----+  
| information_schema |  
| cm |  
| firehose |  
| hue |  
| metastore |  
| mysql |  
| nav |  
| navms |  
| oozie |  
| retail_db |  
| rman |  
| sentry |  
+-----+  
12 rows in set (0.01 sec)  
  
mysql> █
```

```
cloudera@quickstart:~  
File Edit View Search Terminal Help  
Reading table information for completion of table and column names  
You can turn off this feature to get a quicker startup with -A  
  
Database changed  
mysql> show tables;  
+-----+  
| Tables_in_oozie |  
+-----+  
| BUNDLE_ACTIONS |  
| BUNDLE_JOBS    |  
| COORD_ACTIONS  |  
| COORD_JOBS     |  
| OOZIE_SYS      |  
| OPENJPA_SEQUENCE_TABLE |  
| SLA_EVENTS     |  
| SLA_REGISTRATION |  
| SLA_SUMMARY    |  
| VALIDATE_CONN  |  
| WF_ACTIONS     |  
| WF_JOBS        |  
+-----+  
12 rows in set (0.00 sec)  
  
mysql> select *from OOZIE_SYS;  
  
mysql> select *from OOZIE_SYS;  
+-----+-----+  
| name          | data          |  
+-----+-----+  
| db.version    | 3             |  
| oozie.version | 4.1.0-cdh5.13.0 |  
+-----+-----+  
2 rows in set (0.00 sec)  
  
mysql> █
```



```
mysql> create database employees;
Query OK, 1 row affected (0.01 sec)

mysql> show databases;
+-----+
| Database |
+-----+
| information_schema |
| cm |
| employees |
| firehose |
| hue |
| metastore |
| mysql |
| nav |
| navms |
| oozie |
| retail_db |
| rman |
| sentry |
+-----+
13 rows in set (0.00 sec)

mysql> use employees;
Database changed
mysql> create table employee(
-> id int not null primary key,
-> name varchar(20),
-> deg varchar(20),
-> salary int,
-> dept varchar(10));
Query OK, 0 rows affected (0.02 sec)
```

```
mysql> desc employee
-> ;
```

Field	Type	Null	Key	Default	Extra
id	int(11)	NO	PRI	NULL	
name	varchar(20)	YES		NULL	
deg	varchar(20)	YES		NULL	
salary	int(11)	YES		NULL	
dept	varchar(10)	YES		NULL	

5 rows in set (0.00 sec)

```
[cloudera@quickstart ~]$ sqoop import --connect jdbc:mysql://localhost/retail db --password cloudera --username root --table orders --m 1 --target-dir /user/cloudera/anshi
Warning: /usr/lib/sqoop/./accumulo does not exist! Accumulo imports will fail.
Please set $ACCUMULO_HOME to the root of your Accumulo installation.
24/07/29 23:13:22 INFO sqoop.Sqoop: Running Sqoop version: 1.4.6-cdh5.13.0
24/07/29 23:13:22 WARN tool.BaseSqoopTool: Setting your password on the command-line is insecure. Consider using -P instead.
24/07/29 23:13:23 INFO manager.MySQLManager: Preparing to use a MySQL streaming resultset.
24/07/29 23:13:23 INFO tool.CodeGenTool: Beginning code generation
24/07/29 23:13:23 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM `orders` AS t LIMIT 1
24/07/29 23:13:23 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM `orders` AS t LIMIT 1
24/07/29 23:13:23 INFO orm.CompilationManager: HADOOP_MAPRED_HOME is /usr/lib/hadoop-mapreduce
Note: /tmp/sqoop-cloudera/compile/22cd4bd560daff3338da0f157b0365fe/orders.java uses or overrides a deprecated API.
Note: Recompile with -Xlint:deprecation for details.
24/07/29 23:13:25 INFO orm.CompilationManager: Writing jar file: /tmp/sqoop-cloudera/compile/22cd4bd560daff3338da0f157b0365fe/orders.jar
24/07/29 23:13:25 WARN manager.MySQLManager: It looks like you are importing from mysql.
24/07/29 23:13:25 WARN manager.MySQLManager: This transfer can be faster! Use the --direct
24/07/29 23:13:25 WARN manager.MySQLManager: option to exercise a MySQL-specific fast path.
24/07/29 23:13:25 INFO manager.MySQLManager: Setting zero DATETIME behavior to convertToNull (mysql)
24/07/29 23:13:25 INFO mapreduce.ImportJobBase: Beginning import of orders
24/07/29 23:13:25 INFO Configuration.deprecation: mapred.job.tracker is deprecated. Instead, use mapreduce.jobtracker.address
24/07/29 23:13:25 INFO Configuration.deprecation: mapred.jar is deprecated. Instead, use mapreduce.job.jar
24/07/29 23:13:25 INFO Configuration.deprecation: mapred.map.tasks is deprecated. Instead, use mapreduce.job.maps
24/07/29 23:13:26 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032
24/07/29 23:13:26 WARN hdfs.DFSClient: Caught exception
java.lang.InterruptedExcep
    at java.lang.Object.wait(Native Method)
    at java.lang.Thread.join(Thread.java:1281)
    at java.lang.Thread.join(Thread.java:1355)
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DFSOutputStream.java:967)
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.endBlock(DFSOutputStream.java:705)
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStream.java:894)
24/07/29 23:13:26 WARN hdfs.DFSClient: Caught exception
java.lang.InterruptedExcep
    at java.lang.Object.wait(Native Method)
    at java.lang.Thread.join(Thread.java:1281)
    at java.lang.Thread.join(Thread.java:1355)
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DFSOutputStream.java:967)
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.endBlock(DFSOutputStream.java:705)
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStream.java:894)
24/07/29 23:13:26 WARN hdfs.DFSClient: Caught exception
java.lang.InterruptedExcep
    at java.lang.Object.wait(Native Method)
    at java.lang.Thread.join(Thread.java:1281)
    at java.lang.Thread.join(Thread.java:1355)
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DFSOutputStream.java:967)
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeInternal(DFSOutputStream.java:935)
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStream.java:931)
```

## Importing tables in hadoop using sqoop

```

24/07/22 23:23:59 INFO mapreduce.Job: Job job_1721711879175_0002 completed successfully
24/07/22 23:23:59 INFO mapreduce.Job: Counters: 30
  File System Counters
    FILE: Number of bytes read=0
    FILE: Number of bytes written=684668
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=469
    HDFS: Number of bytes written=2999944
    HDFS: Number of read operations=16
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=8
  Job Counters
    Launched map tasks=4
    Other local map tasks=4
    Total time spent by all maps in occupied slots (ms)=61524
    Total time spent by all reduces in occupied slots (ms)=0
    Total time spent by all map tasks (ms)=61524
    Total vcore-milliseconds taken by all map tasks=61524
    Total megabyte-milliseconds taken by all map tasks=63000576
  Map-Reduce Framework
    Map input records=68883
    Map output records=68883
    Input split bytes=469
    Spilled Records=0
    Failed Shuffles=0
    Merged Map outputs=0
    GC time elapsed (ms)=539
    CPU time spent (ms)=5870
    Physical memory (bytes) snapshot=483188736
    Virtual memory (bytes) snapshot=6044938240
    Total committed heap usage (bytes)=243007488
  File Input Format Counters
    Bytes Read=0
  File Output Format Counters
    Bytes Written=2999944
24/07/22 23:23:59 INFO mapreduce.ImportJobBase: Transferred 2.861 MB in 29.4574 seconds (99.4533 KB/sec)
24/07/22 23:23:59 INFO mapreduce.ImportJobBase: Retrieved 68883 records.
[cloudera@quickstart ~]$ sqoop import --connect jdbc:mysql://quickstart:3306/retail_db --password cloudera --username root --table orders;

```

To check the imported tables

Hadoop fs -ls

```

drwxr-xr-x - cloudera cloudera 0 2024-07-22 23:19 departments
drwxr-xr-x - cloudera cloudera 0 2024-07-22 23:23 orders

```

To make a table in sql where the data will be exported

```

[cloudera@quickstart ~]$ mysql -u root -pcloudera
Welcome to the MySQL monitor.  Commands end with ; or \g.
Your MySQL connection id is 58
Server version: 5.1.73 Source distribution

```

Copyright (c) 2000, 2013, Oracle and/or its affiliates. All rights reserved.

Oracle is a registered trademark of Oracle Corporation and/or its affiliates. Other names may be trademarks of their respective owners.

Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.

```

mysql> use retail_db;
Reading table information for completion of table and column names
You can turn off this feature to get a quicker startup with -A

```

```

Database changed
mysql> select * from empty;
Empty set (0.00 sec)

```

```
[cloudera@quickstart ~]$ sqoop export --connect jdbc:mysql://localhost/retail_db --username root --password cloudera --table empty --export-dir /user/cloudera/toexport/data.txt
```

To export the table from local to cloud to sql

To check the exported table

```
mysql> use retail_db;  
Reading table information for completion of table and column names  
You can turn off this feature to get a quicker startup with -A
```

Database changed

```
mysql> select * from empty;
```

```
+-----+  
| id | name   | location |  
+-----+  
| 1  | Ruchit | dombivli |  
| 2  | Prisha | mulund   |  
| 3  | Khushi | kandivali |  
| 4  | Kaushal | Borivali |  
+-----+  
4 rows in set (0.00 sec)
```

```
mysql>
```

---